

# COMPX216/Y05337

# Artificial Intelligence

## Problem-solving agents

# Today: Problem-solving agents

- The agent and the environment
- Types of environments
- Problem-solving agents
- The problem-solving process
- Components of a search problem
- Example problems
- Search trees and search data structures
- Expanding nodes in a search tree
- The breadth-first search algorithm

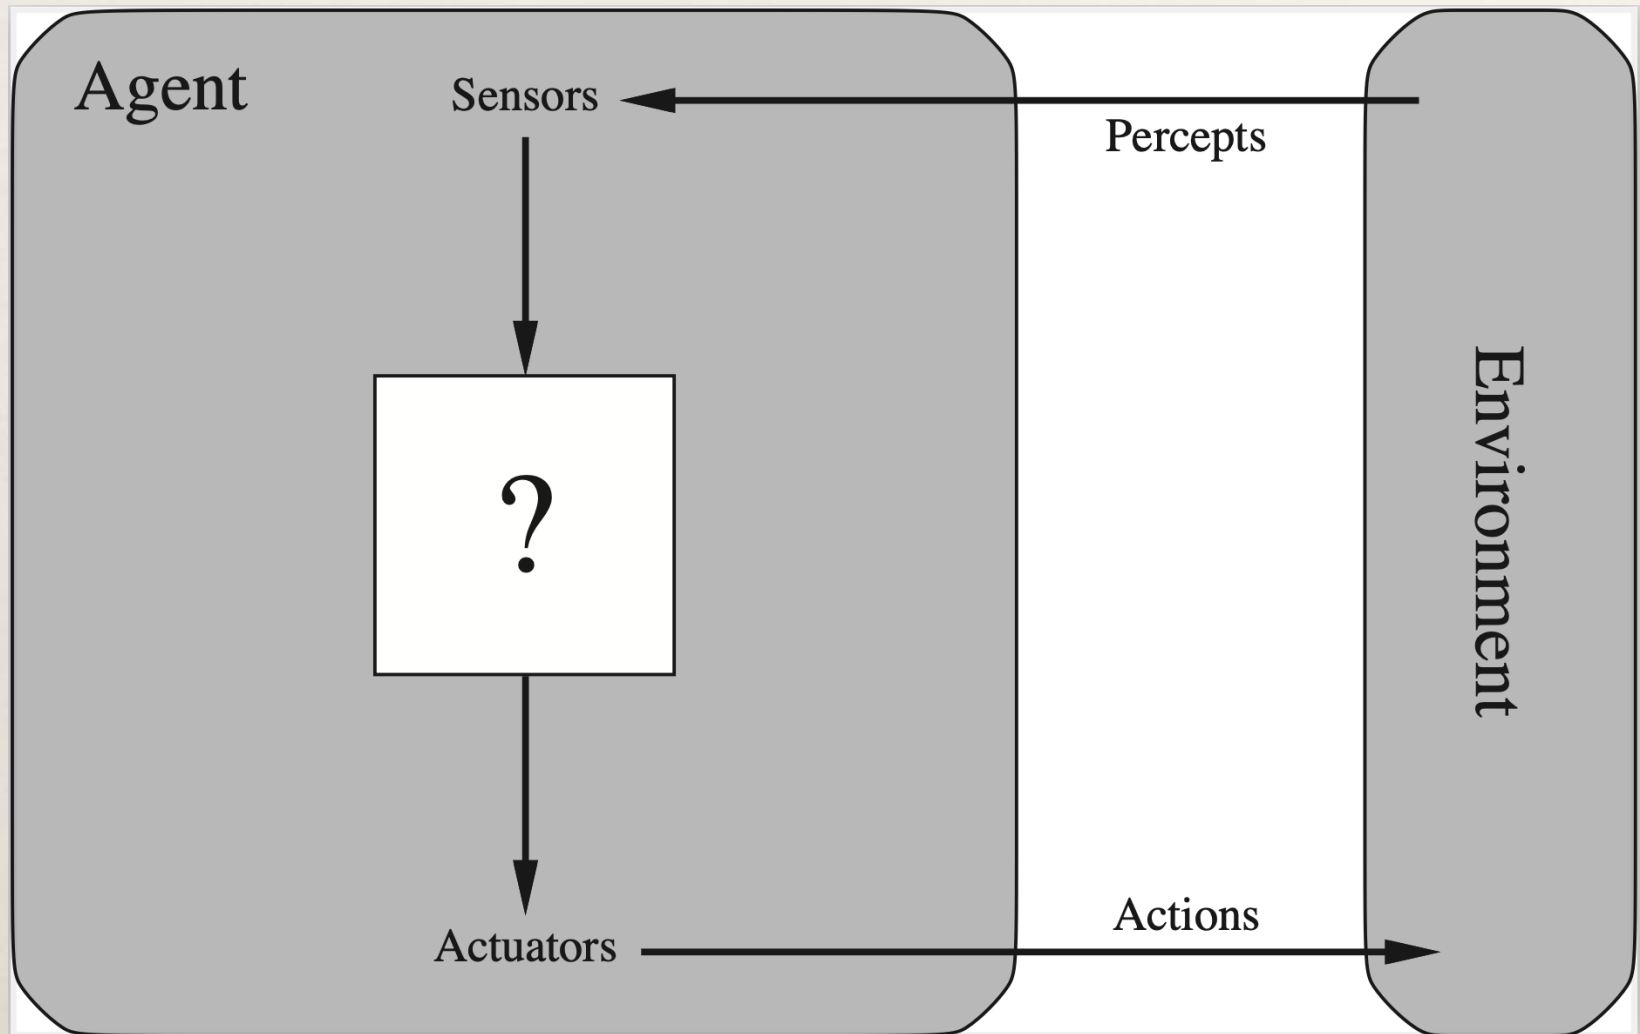
# What is an Agent

- An agent exists in a specific environment
- The agent is the entity for which we want to make decisions and take actions
- In most cases, the environment defines a problem, while the agent gives solution

# Agents and Environments

- An *agent* within an *environment* will take *actions* to achieve a particular *goal*
- The agent *perceives* its environment somehow, and makes a *decision* about which action to take
- The agent takes its action(s), which *changes* the environment, and the process *repeats*
- Artificial Intelligence is the algorithmic discipline of making these decisions more intelligent

# The agent and the environment



# Agents

- Perceives its environment through sensors
  - Robot: Camera, Laser, Sonar
  - Human: 5 Senses
  - Games: Keyboard Input / RAM / Game State
- Acts on its environment through Actuators
  - Robots: Wheels, Tracks, Arms
  - Humans: Limbs, Tools
  - Games: Predefined Rules / Actions
- May have some knowledge about environment

# Agents

- Percept
  - Agent's perceptual inputs at any instant
  - "Agent's current belief of the world"
  - Environment's current State
- Percept Sequence
  - Complete history of everything the agent has ever perceived so far

# Agents

- *In general*, an agent's choice of action at any given instant can depend on the entire percept sequence to date (what you did in past matters)
- Agent's behavior is given by the agent function that maps given percept sequence to an action
  - "Given what I have seen, take this action"



# Agents


- An action taken by an agent in a given state transitions the state to another
- State Transition Function (STF)
  - Successor Function, Table, Graph
- State  $S$ , Action  $A$
- $S' = \text{STF}(S, A)$

# Agents


**State Transition  
Table**

Input State	1	0
$S_1$	$S_1$	$S_2$
$S_2$	$S_2$	$S_1$

**State Diagram**



# A very simple vacuum cleaner agent



## Percept

[A, Clean]

[A, Dirty]

[B, Clean]

[B, Dirty]

## Action

Move right



Suck dirt

Move left

Suck dirt

# Types of environments

- Fully observable vs. partially observable
  - Do the sensors provide access to the complete state of the environment?



# Types of environments

- **Single-agent vs. multi-agent**
  - Note that multi-agent environments can be competitive or cooperative





# Types of environments

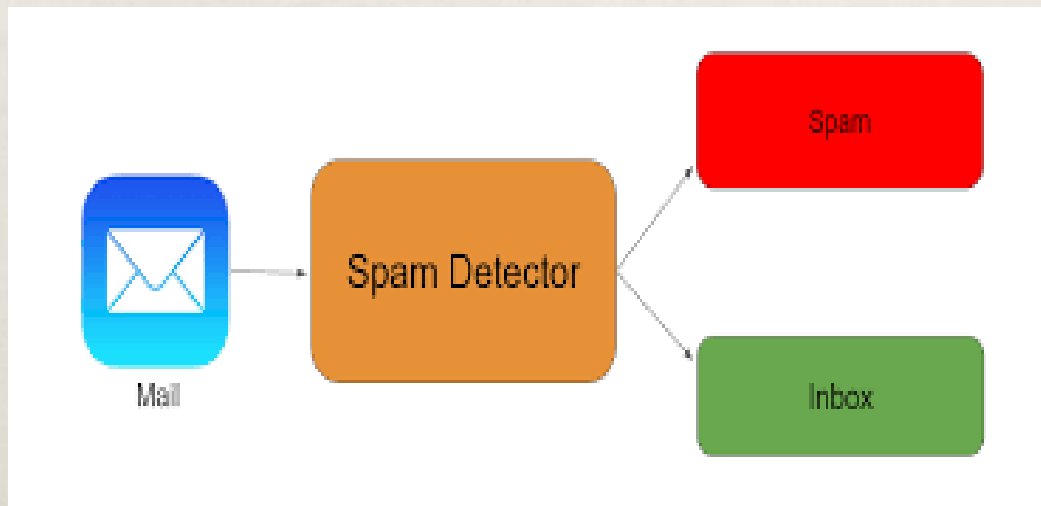
- **Deterministic vs. non-deterministic**
  - Note that partially observable environments may *appear* nondeterministic



# Types of environments

- Episodic vs. sequential

- Does the outcome of an action only depend on the current state?



# Types of environments

- Static vs. dynamic
  - Can the environment change while the agent is choosing an action?

## EASY SUDOKU

8		1	3	4			2	
	5		6			8		3
				9	5	1		
6				5	9			4
		3				7	5	
		5	2	3		6	8	
		9	5		8	4		6
5	7		1			2		8
3		6						

PUZZLES FOR ADULTS






# Types of environments

- Discrete vs. continuous


- This distinction can apply to actions, percepts, states, and time



# Problem-solving agents

- Correct action may not be obvious: we need to find a sequence of actions that form a **path** to the **goal state**
- We can do this by performing **search** in the **state space**
- **Problem-solving agents** use **atomic** state representations: the state of the world is considered indivisible
  - More complex **planning agents** are used for problems with **structured** or **factored** state representations
- The environments we consider are fully observable, single-agent, deterministic, episodic, static, and discrete
- We also assume that the effects of actions on the environment are known (i.e., its “physics” are known)

# State-space graph for vacuum cleaner problem



# The problem-solving process

- We can use the following four-step process to apply problem-solving agents to real-world problems
- **Goal formulation:** what is the goal of the search?
- **Problem formulation:** what is a suitable, sufficiently abstract model of relevant states and actions?
- **Search:** identify a sequence of actions that reach a goal state in the model, i.e., a path that is a **solution**
- **Execution:** perform the sequence of actions in the actual real-world environment
- We will look at algorithms for the third step

# Components of a search problem

- **States:** the set of states the environment can be in
- **Initial state:** the starting point for the search
- **Goal states:** this set is implicitly defined by a Boolean function that tests whether a state is a goal state
- **Actions:** a function that returns the finite set of actions that are applicable in a given state
- **Transition model:** a function that defines the next state for a given state-action pair
- **Action cost function:** a function that defines a numeric cost of performing an action to move between two states
  - **Optimum solutions** are those with the lowest **path cost**



# The vacuum world: problem definition

- **States:** a grid of cells, where each cell is either dirty or not, and the robot is in exactly one cell
- **Initial state:** any state in the state space
- **Goal states:** states in which all cells are clean
- **Actions:** *suck, move left, move right* (in a 2D world, we could add *move up* and *move down*)
- **Transition model:** a function that implements the transitions in the state-space graph from a few slides back
- **Action cost function:** each action has a cost of 1
  - This means the past cost is the number of actions it involves

# The sliding-tile puzzle

- Most well-known sliding-tile puzzle: 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

# The sliding-tile puzzle: problem definition

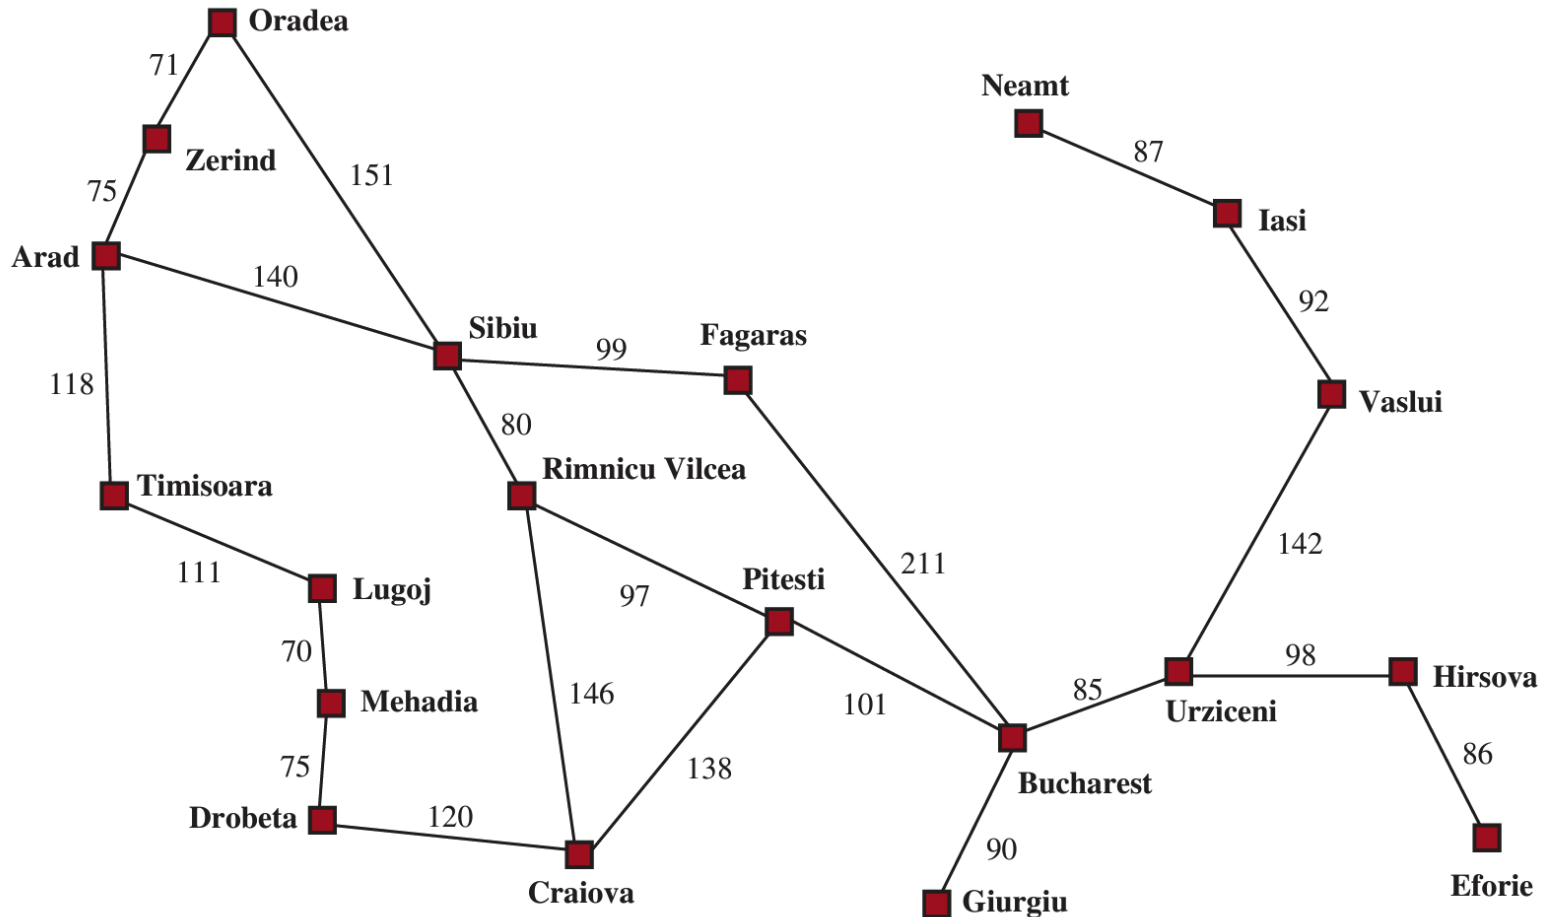
- **States:** these specify the locations of all the tiles
- **Initial state:** any state in the state space
- **Goal states:** states where the numbers are in order
- **Actions:** moving the blank space *up, down, left, or right*, considering only applicable actions
- **Transition model:** a function that swaps the blank with an adjacent number based on the given action
- **Action cost function:** each action has a cost of 1



# Knuth's conjecture & problem definition

- We can get any positive integer by starting with 4 and applying the square root, the floor function, or the factorial
- **States:** positive real numbers (note that this is infinite!)
- **Initial state:** the number 4
- **Goal state:** the desired positive integer
- **Actions:** apply the square root, floor function, or factorial (if applicable) to the current state
- **Transition model:** given by the definitions of the mathematical functions that are available
- **Action cost function:** each action has a cost of 1

# Route finding (in Romania)




# Route finding: problem definition

- **States:** locations on the map
- **Initial state:** any chosen location
- **Goal state:** the desired destination
- **Actions:** moves to adjacent cities on the map
- **Transition model:** given by the graph representing the map
- **Action cost function:** each action has a cost given by the distance between the two corresponding cities
- The optimum solution is the path with the shortest distance between the initial state and the goal state

# Search trees: exploring paths in the state space

- Search algorithms construct a search tree that has nodes representing states on the paths generated so far



# Search data structures

- Node objects represent nodes using five instance attributes:
  - `self.state`: a reference to an object representing a state
  - `self.parent`: a reference to the node in the tree from which this node was generated by performing an action
  - `self.action`: the action that was performed on the parent node's state to generate this node
  - `self.path_cost`: the total cost of the path from the initial state to this node
  - `self.depth`: the depth of the node in the search tree
- We also need to keep track of the **generated** but not yet **expanded** nodes (the so-called search **frontier**) using a **queue**
- Additionally, we should use a hashtable to check for states that have been **reached** (i.e., generated) previously

# Node vs State

- State
  - Configuration of the environment
- Node
  - Bookkeeping data structure
  - Exists only within the search tree
  - Nodes are "on path" in the search tree



# Expanding nodes

- To generate the search tree, we need a function that takes a node and uses the problem definition to generate children

```
function EXPAND(problem, node) yields nodes
    s ← node.STATE
    for each action in problem.ACTIONS(s) do
        s' ← problem.RESULT(s, action)
        cost ← node.PATH-COST + problem.ACTION-COST(s, action, s')
        yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```

- Note that this uses the yield keyword to indicate a resumable function (see generator functions Python!)

# Search strategies

- Uninformed (blind) search
  - No information about states beyond problem description
  - Limited to children generation, goal test
- Informed (heuristic) search
  - Can guess which states are 'more promising'
  - Hopefully leads to faster search episodes



# Which strategy to choose?

- Completeness
  - Is it guaranteed to find a solution if it exists?
- Optimality
  - Does it find the optimal solution?
- Time complexity
  - How long does it take to find a solution?
- Space complexity
  - How much memory is need to run the search?