

# COMPX216-24A

# Artificial Intelligence

Local search and optimisation


# Today: Local search and optimisation

- Optimisation problems
- Local search
- Hill-climbing search
- Local and global optima
- Ridges and plateaus
- The 8-queens problem
- Variants of hill-climbing
- Simulated annealing
- Local beam search

# Optimisation problems

- In many applications, we just want to find a good state, not a sequence of actions that get us to a goal from an initial state
- Example: given items of different weight, how can we distribute them as fairly as possible into two backpacks?
  - This is an instance of the *number partitioning problem*
- Finding a good state is called an **optimisation problem**
- In the context of optimization problems, the evaluation function used for states is called the **objective function**
- Depending on the problem, we may want to find a state where the objective function is **maximised** or **minimised**

# Recall: Search problems




# Local search problems

											<b>B</b>

# Local search


- Optimisation problems are often tackled using **local search** algorithms that may not find the best possible state
- Local search algorithms do not keep track of paths used to reach states and do not keep track of states reached previously
- In the simplest case, they move around the state space by keeping information only about the state that looks best so far
- Local search algorithms require very limited memory but are unsystematic in their search

# Local search and optimisation



# Local search and optimisation


Cost: 17





# Local search and optimisation

Cost: 17

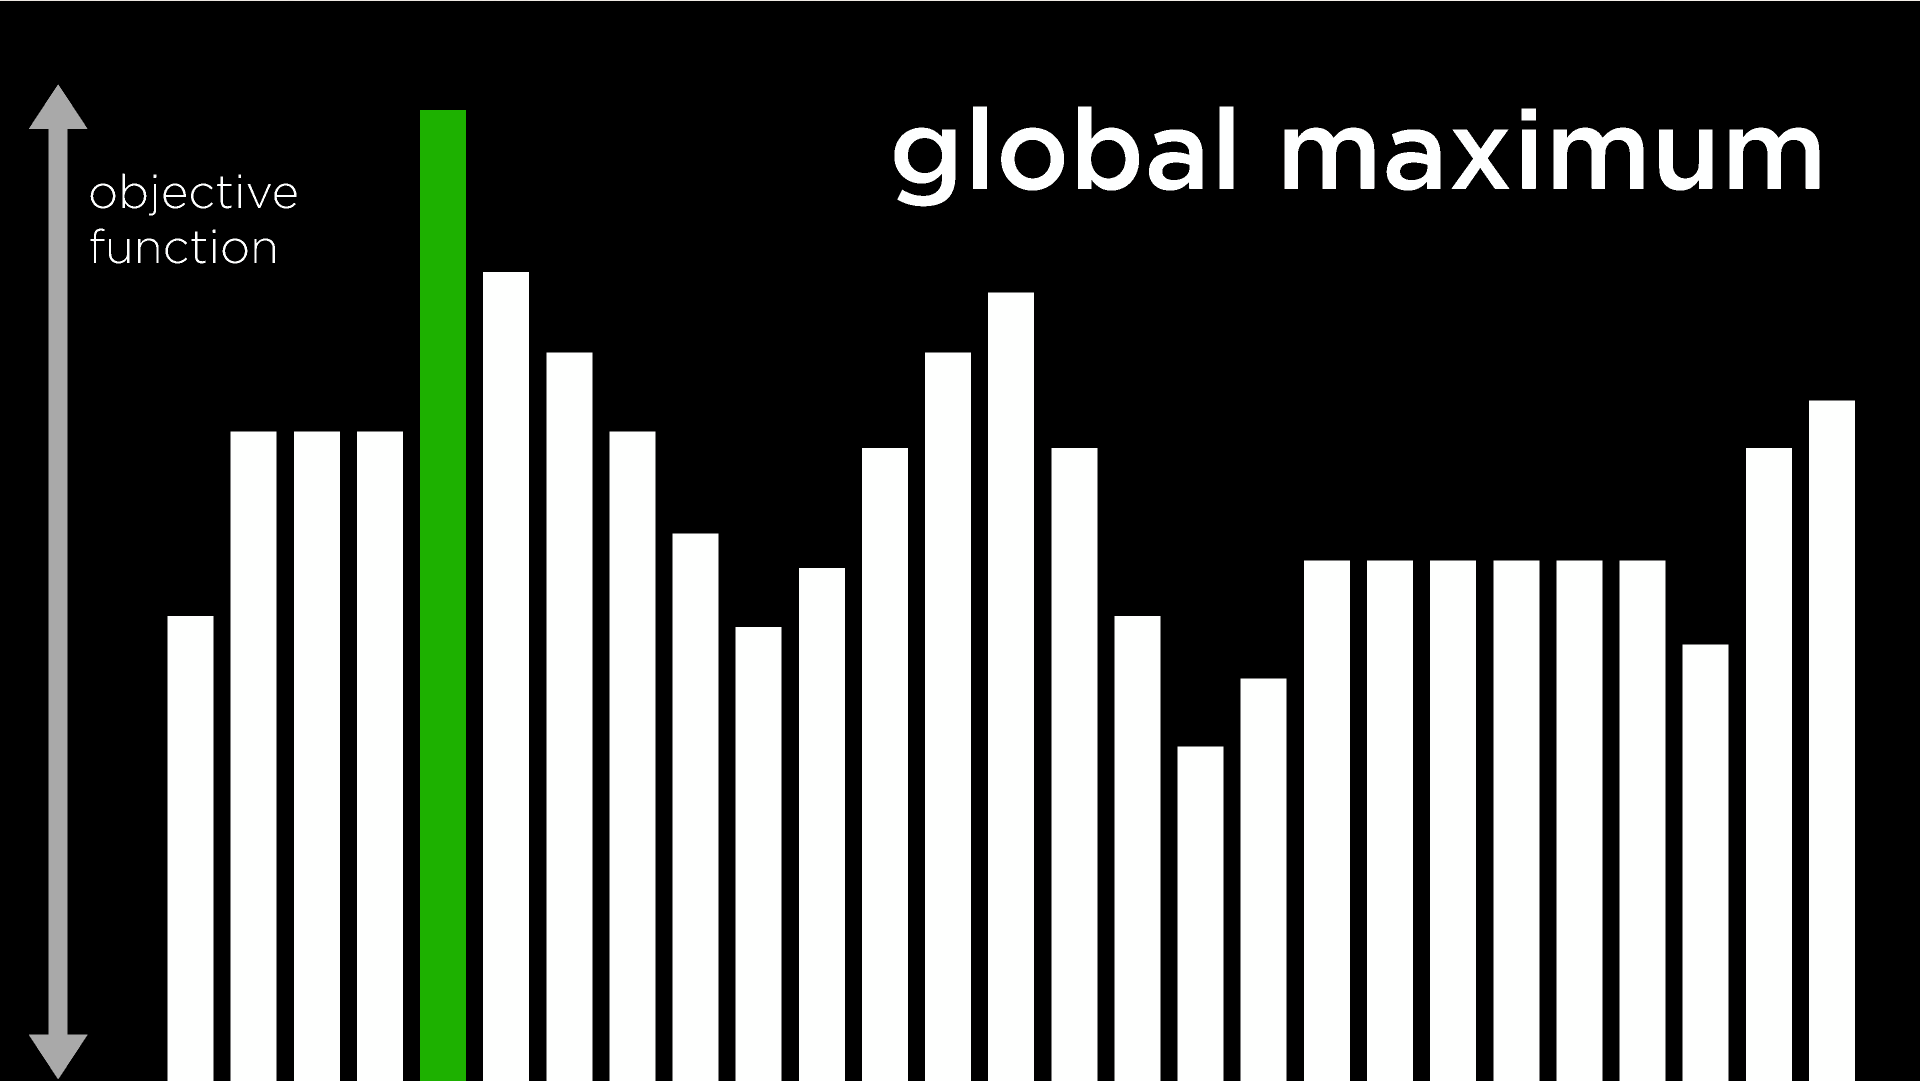


# Local search and optimisation


state-space landscape



# Local search and optimisation




# Local search and optimisation




# Local search and optimisation

current state



# Local search and optimisation

neighbors




# Hill-climbing search

- Hill-climbing search is a type of optimisation problem.
- Moves to the best neighbour in state space in each iteration until no neighbour is better
- More sophisticated variants keep information about a fixed-size set of states




# Hill-climbing search







# Hill-climbing search




# Hill-climbing search




# Hill-climbing search




# Hill-climbing search




# Hill-climbing search




# Hill-climbing search



# Hill-climbing search




# Hill-climbing search






# Hill-climbing search



# Hill-climbing search




# Hill-climbing search



```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  current  $\leftarrow$  problem.INITIAL
  while true do
    neighbor  $\leftarrow$  a highest-valued successor state of current
    if VALUE(neighbor)  $\leq$  VALUE(current) then return current
    current  $\leftarrow$  neighbor
```

# The $n$ -queens problem



- The objective is to minimise the number of pairs of queens that attack each other
- Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal
- Move a queen to reduce number of conflicts



$h = 5$



$h = 2$




$h = 0$

- Almost always solves  $n$ -queens problems almost instantaneously for very large  $n$ , e.g.,  $n = 1 \text{ million}$


# Hill-climbing search and optimisation

Cost: 17




# Hill-climbing search and optimisation

Cost: 17




# Hill-climbing search and optimisation

Cost: 17



# Hill-climbing search and optimisation


Cost: 15






# Hill-climbing search and optimisation

Cost: 13




# Hill-climbing search and optimisation

Cost: 11




# Hill-climbing search and optimisation

Cost: 9



# Hill-climbing search and optimisation




# Hill-climbing search and optimisation

global maximum



# Hill-climbing search and optimisation

local maxima




# Hill-climbing search and optimisation

global minimum




# Hill-climbing search and optimisation

local minima







# Hill-climbing search and optimisation




# Hill-climbing search and optimisation




# Hill-climbing search and optimisation



# Hill-climbing search and optimisation




# Hill-climbing search and optimisation



# Hill-climbing search and optimisation

flat local maximum




# Hill-climbing search and optimisation

shoulder



# Hill-climbing search






# Variants of hill-climbing

- **Stochastic hill-climbing** chooses at random from uphill moves
  - Probability of selection may be based on size of improvement
- Stochastic hill-climbing finds good states more slowly but may find better solutions in some state spaces
- **First-choice hill-climbing** generates neighbours randomly and picks the first one that yields an improvement
  - Useful if the set of neighbours is very large
- Good, very simple option to improve on basic hill-climbing: **random-restart hill-climbing**
  - If each of  $n$  restarts has probability  $p$  of not finding the global maximum, the probability of failure is  $1 - p^n$


# Simulated annealing

- In metallurgy, **annealing** is used to increase “ductility” and decrease hardness of a material to make it easier to work with
- Hot material’s temperature is gradually lowered according to a pre-defined schedule to achieve this
- **Simulated annealing**: stochastic hill-climbing that allows downhill moves with a probability dependent on the temperature
- Early on, higher "temperature": more likely to accept neighbors that are worse than current state
- Later on, lower "temperature": less likely to accept neighbors that are worse than current state


# Simulated annealing




# Simulated annealing




# Simulated annealing




# Simulated annealing



# Simulated annealing




# Simulated annealing







# Simulated annealing



# Simulated annealing



# Simulated annealing



# Simulated annealing

```
function SIMULATED-ANNEALING(problem, max):  
    current = initial state of problem  
    for t = 1 to max:  
        T = TEMPERATURE(t)  
        neighbor = random neighbor of current  
         $\Delta E$  = how much better neighbor is than current  
        if  $\Delta E > 0$ :  
            current = neighbor  
        else  
            with probability  $e^{\Delta E/T}$  set current=neighbor  
    return current
```

# Temperature analogies

## Searching for a New Apartment:

At the start, when looking for a **new apartment**, you might visit many different neighborhoods and consider various layouts (**high temperature**). Over time, you **narrow down** to a few best choices and finally commit to one (**low temperature**).

- **High temperature:** Willing to explore various locations, even ones that seem inconvenient.
- **Low temperature:** Focusing only on the best options and making a final choice.



<https://propertyrecordsofmaryland.com/how-to-find-property-owner-information-in-maryland-tips-and-tools/>

COMPX216-24A

# Temperature analogies

## Decision-Making When Shopping

When you first enter a **shopping mall**, you're open to exploring multiple stores and trying out different products (**high temperature, more randomness**). As time passes and you get tired, you focus on finalizing a purchase (**low temperature, less randomness**), settling on the best option you've found.

- **Early stage (high temp):** Exploring many product options, including ones that might not seem great at first.
- **Later stage (low temp):** Narrowing down choices and making a final decision.



<https://www.campaignasia.com/article/economic-hurdles-push-japanese-consumers-to-save-not-spend/452840>

# Local beam search

- The local search algorithms so far kept track of a single state
- **Local beam search** keeps track of  $k$  states instead
- In each iteration, all the successors of all  $k$  states are generated
- Unless a stopping condition is met, the best  $k$  successors are selected for the next iteration
  - Possible stopping condition: no improvement found
- This seems similar to random restart hill-climbing but is actually different: hill-climbing runs are executed *independently*
- Local beam search automatically concentrates the search on those parts of the explored search space where progress is most rapid
- We can increase exploration by adopting **stochastic beam search**
  - Chooses successors with probability proportional to their value

# References

- CS50's Introduction to Artificial Intelligence with Python 2020  
(<https://cs50.harvard.edu/ai/2024/>)